


```
In [113]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [114]: from keras.datasets import mnist
```

```
In [115]: (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
In [116]: plt.imshow(X_train[1])
```

Out[116]: <matplotlib.image.AxesImage at 0x181178ed908>



```
In [117]: print(X_train[1].max(), X_train[1].min())
```

255 0

```
In [118]: y_train
```

Out[118]: array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)

```
In [119]: X_train = X_train/255
```

```
In [120]: X_test = X_test/255
```

```
In [121]: from keras.utils import to_categorical
```

```
In [122]: y_cat_train = to_categorical(y_train)
```

```
In [123]: y_cat_train
```

```
Out[123]: array([[0., 0., 0., ..., 0., 0., 0.],
                  [1., 0., 0., ..., 0., 0., 0.],
                  [0., 0., 0., ..., 0., 0., 0.],
                  ...,
                  [0., 0., 0., ..., 0., 0., 0.],
                  [0., 0., 0., ..., 0., 0., 0.],
                  [0., 0., 0., ..., 0., 1., 0.]], dtype=float32)
```

```
In [124]: y_cat_test = to_categorical(y_test)
```

```
In [125]: X_train = X_train.reshape(60000, 28, 28, 1)
```

```
In [126]: X_test = X_test.reshape(10000, 28, 28, 1)
```

```
In [127]: from keras.models import Sequential
          from keras.layers import Dense, Conv2D, MaxPool2D, Flatten, Dropout
          from keras.callbacks import EarlyStopping
```

```
In [128]: model = Sequential()
```

```
In [129]: model.add(Conv2D(filters=32, kernel_size=(4,4), activation='relu', padding='same', input_shape=(28,28,1)))
          model.add(MaxPool2D())
          model.add(Flatten())
          model.add(Dense(units=64, activation='relu'))
          model.add(Dense(units=32, activation='relu'))
          model.add(Dense(units=10, activation='softmax'))
          model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [132]: early_stop = EarlyStopping(monitor='val_loss', patience=4)
```

```
In [133]: model.fit(X_train, y_cat_train, batch_size=200, epochs=15, validation_data=(X_test, y_cat_test), callbacks=[early_stop])
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/15

60000/60000 [=====] - 71s 1ms/step - loss: 0.2985 - accuracy: 0.9143 - val_loss: 0.0896 - val_accuracy: 0.9732

Epoch 2/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0840 - accuracy: 0.9747 - val_loss: 0.0655 - val_accuracy: 0.9789

Epoch 3/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0565 - accuracy: 0.9825 - val_loss: 0.0501 - val_accuracy: 0.9836

Epoch 4/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0431 - accuracy: 0.9866 - val_loss: 0.0441 - val_accuracy: 0.9846

Epoch 5/15

60000/60000 [=====] - 70s 1ms/step - loss: 0.0347 - accuracy: 0.9896 - val_loss: 0.0420 - val_accuracy: 0.9856

Epoch 6/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0286 - accuracy: 0.9911 - val_loss: 0.0431 - val_accuracy: 0.9863

Epoch 7/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0226 - accuracy: 0.9930 - val_loss: 0.0459 - val_accuracy: 0.9850

Epoch 8/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0180 - accuracy: 0.9946 - val_loss: 0.0528 - val_accuracy: 0.9839

Epoch 9/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0156 - accuracy: 0.9954 - val_loss: 0.0407 - val_accuracy: 0.9862

Epoch 10/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0119 - accuracy: 0.9964 - val_loss: 0.0427 - val_accuracy: 0.9869

Epoch 11/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0095 - accuracy: 0.9973 - val_loss: 0.0418 - val_accuracy: 0.9869

Epoch 12/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0087 - accuracy: 0.9974 - val_loss: 0.0540 - val_accuracy: 0.9843

Epoch 13/15

60000/60000 [=====] - 69s 1ms/step - loss: 0.0079 - accuracy: 0.9976 - val_loss: 0.0501 - val_accuracy: 0.9861

```
Out[133]: <keras.callbacks.callbacks.History at 0x1811977a448>
```

```
In [134]: model_hist = pd.DataFrame(model.history.history)
```


```
In [135]: model_hist.head()
```

```
Out[135]:
```

	val_loss	val_accuracy	loss	accuracy
0	0.089607	0.9732	0.298495	0.914350
1	0.065522	0.9789	0.084036	0.974717
2	0.050092	0.9836	0.056484	0.982467
3	0.044078	0.9846	0.043068	0.986633
4	0.042035	0.9856	0.034715	0.989567


```
In [136]: model_hist[['val_accuracy', 'accuracy']].plot()
```

```
Out[136]: <matplotlib.axes._subplots.AxesSubplot at 0x18119286f88>
```



```
In [137]: model_hist[['val_loss', 'loss']].plot()
```

```
Out[137]: <matplotlib.axes._subplots.AxesSubplot at 0x18142ed37c8>
```



```
In [138]: y_pred = model.predict_classes(X_test)
```

```
In [139]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

```
In [141]: accuracy_score(y_test, y_pred)
```

```
Out[141]: 0.9861
```

```
In [142]: confusion_matrix(y_test, y_pred)
```

```
Out[142]: array([[ 977,    0,    1,    1,    0,    0,    1,    0,    0,    0],
 [    0, 1125,    3,    2,    0,    0,    2,    1,    2,    0],
 [    0,    1, 1023,    0,    1,    0,    0,    7,    0,    0],
 [    0,    0,    4, 1002,    0,    3,    0,    1,    0,    0],
 [    2,    1,    6,    0,  948,    0,    3,    1,    3,   18],
 [    2,    0,    0,    8,    0,  880,    1,    0,    0,    1],
 [    5,    2,    1,    0,    1,    1,  947,    0,    1,    0],
 [    0,    0,    7,    0,    0,    0,    0, 1015,    2,    4],
 [    5,    0,    7,    2,    0,    2,    1,    1,  951,    5],
 [    0,    1,    0,    4,    3,    2,    0,    5,    1,  993]],
      dtype=int64)
```

```
In [144]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	980
1	1.00	0.99	0.99	1135
2	0.97	0.99	0.98	1032
3	0.98	0.99	0.99	1010
4	0.99	0.97	0.98	982
5	0.99	0.99	0.99	892
6	0.99	0.99	0.99	958
7	0.98	0.99	0.99	1028
8	0.99	0.98	0.98	974
9	0.97	0.98	0.98	1009
accuracy			0.99	10000
macro avg	0.99	0.99	0.99	10000
weighted avg	0.99	0.99	0.99	10000

```
In [ ]:
```